

VIKAS PASPULA

AI Engineer — Full-Stack Developer (LLMs • RAG • Cloud-Native Systems)
Hyderabad, Telangana
vikaspaspula9@gmail.com — 9100539058 — LinkedIn — GitHub

PROFESSIONAL SUMMARY

AI Engineer with hands-on experience in designing, fine-tuning, and deploying LLM-powered systems, RAG pipelines, and cloud-native full-stack applications. Proven ability to integrate NLP models into production environments using Python, PyTorch, LangChain, Docker, and CI/CD pipelines. Experienced in scalable AI microservices, semantic search, and enterprise-grade deployment across education, HR tech, and conversational AI domains.

EDUCATION

Mahindra University <i>Master of Engineering in Artificial Intelligence and Data Science</i> <i>Aug 2019 – Jun 2021</i>	Hyderabad , Telangana <i>CGPA: 7.60%</i>
Sreenidhi Institute of Science and Technology <i>Bachelor of Engineering in Computer and Electronics Engineering;</i> <i>Jun 2015 – Jun 2019</i>	Hyderabad,Telangana <i>CGPA: 7.24%</i>
Narayana Junior College <i>July 2013 – May 2015</i>	Senior Secondary (XII); <i>CGPA: 8.63%</i>

EXPERIENCE

Gen AI Developer <i>Chandigarh</i>	Cipher Schools <i>Jun 2023 – Present</i>
<ul style="list-style-type: none">Designed and fine-tuned lightweight LLMs (CodeLLaMA, StarCoder) to detect AI-generated code using semantic embeddings and stylistic analysis.Built classifiers combining behavioral metadata and transformer embeddings, outperforming traditional ML approaches in code authorship detection.Optimized LLM inference pipelines for scalability and cost efficiency in production environments.	
AI Engineer <i>Hyderabad</i>	Tech Mahindra <i>Nov 2021 – Jun 2023</i>
<ul style="list-style-type: none">Developed an AI-driven resume screening system that automated 80% of manual candidate matching using NLP and semantic similarity.Implemented transformer-based QA systems using dense vector search for enterprise knowledge retrieval.Designed ethical, multi-turn conversational workflows with explainability and prompt-safety constraints.	
AI Engineer <i>Hyderabad, Telangana</i>	Ineuron <i>Feb 2019 – Oct 2021</i>
<ul style="list-style-type: none">Built end-to-end AI/ML applications covering supervised and unsupervised learning pipelines.Developed RESTful backend services with PostgreSQL to support multi-user data access and scalable transactions.Collaborated on deploying ML models into production-ready web applications.	

PROJECTS

Enterprise AI Job Portal with Resume Intelligence
<ul style="list-style-type: none">Built a full-stack job portal using React, Spring Boot, and PostgreSQL with authentication and role-based access.Integrated NLP-based resume parsing to extract 10–15 key skills per profile, automating candidate-job matching.Deployed frontend on Vercel and backend on Render with Dockerized services and CI/CD workflows.

Cloud-Native AutoML Platform

- Designed a scalable AutoML pipeline supporting PostgreSQL, MySQL, and MongoDB data sources.
- Implemented preprocessing, model evaluation, and reporting with accuracy and confusion-matrix metrics.

Context-Aware Conversational AI System (RAG)

- Built an LLM-powered chatbot using LangChain, FAISS, and ChromaDB for contextual retrieval.
- Designed secure, multi-turn conversational flows with explainable AI responses.

Photo-Realistic Image Super-Resolution (SRGAN)

- Implemented SRGAN using PyTorch to enhance low-resolution images.
- Evaluated model performance using PSNR and SSIM, with applications in medical imaging and surveillance.

SKILLS

Programming: Python, Java, C++, OOPs, DSA

AI Agent Development: Agentic AI, RAG, LangChain, Tool-Calling Agents, Vector Indexing (FAISS/ChromaDB), Conversational AI, Multi-Turn Dialogue

Google AI Ecosystem: Vertex AI (AutoML-style workflows, RAG, Evaluation), Gemini Enterprise-style conversational orchestration, Cloud deployment patterns for AI workloads

Natural Language Processing: BERT, LSTM, GRU, Transformers, Attention, RAG

Backend: FastAPI/Django/Spring Boot, REST APIs, JWT Authentication

MLOps / ML frameworks: PyTorch, TensorFlow, Scikit-learn

Databases: PostgreSQL, MySQL, MongoDB

DevOps / Tools: Docker, CI/CD, Git/GitHub

Languages: Telugu (Native), English (Professional), Hindi